

Machine Learning Assignment No: 2 (Classification).

Name: Khuram Shahzad, Roll No: p218742

ASSIGNMENT QUESTION:

Classify the following dataset to distinguish between skin color and nonskin color. You may try different models but report at least the two top scoring classifiers. Use the following settings for the experiments.

1. Stratified 5-Fold cross-validation (the dataset is imbalanced)
2. Report precision, recall, and F1-score (with micro as averaging scheme)

Data Set: <https://archive.ics.uci.edu/ml/datasets/Skin+Segmentation>
(<https://archive.ics.uci.edu/ml/datasets/Skin+Segmentation>).

Submit a code document and a word document with setup, results and discussion.

Importing Libraries

```
In [14]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix#for visualization
from sklearn.tree import plot_tree
from sklearn.metrics import precision_recall_curve, f1_score
from sklearn.model_selection import StratifiedKFold
from statistics import mean, stdev
from sklearn import datasets
from sklearn.metrics import confusion_matrix, roc_auc_score ,roc_curve,auc
import seaborn as sns
from sklearn import metrics
from sklearn.tree import DecisionTreeClassifier#for checking testing results
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
```

Importing the Data

```
In [15]: def load_data():
    '''
    Loading the data for this file into a pandas DataFrame.
    '''
    frame = pd.read_csv(
        "Skin_NonSkin.txt",

        # Specify the file encoding
        encoding='utf-8', # UTF-8 is common

        # Specify the separator in the data
        sep='\t',          # tab separated values

        # Ignore spaces after the separator
        skipinitialspace=True,

        # Generate row labels from each row number
        index_col=None,

        # Generate column headers row from each column number
        header=0,
    )
    # Return the entire frame
    return frame
```

Exploratory Data Analysis (EDA)

```
In [13]: #reading the data
def ExploratoryDataAnalysis(i):
    df= load_data()
    df.head()
    #getting information of dataset
    df.info()
    df.head()
    if i==1:
        # let's plot pair plot to visualise the attributes all at once
        sns.pairplot(data=load_data(), hue = 'L1')
    ExploratoryDataAnalysis(0)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 245057 entries, 0 to 245056
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0    c1      245057 non-null    int64
1    c2      245057 non-null    int64
2    c3      245057 non-null    int64
3    L1      245057 non-null    int64
dtypes: int64(4)
memory usage: 7.5 MB
```


We understand that this dataset has 245057 records, 4 columns with the type int and there are no

NAN values as form following command

Now we perform some basic EDA on this dataset. Let's check the correlation of all the features with each other

```
In [22]: # Let's plot pair plot to visualise the attributes all at once
sns.pairplot(data=load_data(), hue = 'L1')
```

```
Out[22]: <seaborn.axisgrid.PairGrid at 0x2656a530dc0>
```



We have a total of 2 targets that we want to predict: 1, and 2. We can see that 1 always forms a different cluster from the 2.

Data Preprocessing

Now, we will separate the target variable(y) and features(X) as follows

```
In [12]: def DataPreprocessing():
df = load_data()
target = df['L1']
df1 = df.copy()
df1 = df1.drop('L1', axis =1)
# Defining the attributes
X = np.delete(df1.to_numpy(), (0), axis=0)
y = np.delete(target.to_numpy(), (0), axis=0)
# converting target values to (0, 1)
t=0;
for i in y:
    if i==1:

        y[t]= 0
    else:
        y[t]= 1
    t=t+1
return X, y
```

Training and Testing

Stratified 5-Fold cross-validation (the dataset is imbalanced)

```

In [11]: def ApplyClassifier(model):

    model_name = type(model).__name__
    kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=1)
    model_accu_stratified = []
    cv_score = []
    i=1

    print('-----')
    print('Classifier Name: ', model_name)
    print('-----')
    for train_ix, test_ix in kfold.split(X, y):
        # select rows
        train_X, test_X = X[train_ix], X[test_ix]
        train_y, test_y = y[train_ix], y[test_ix]

        # summarize train and test composition
        train_1, train_2 = len(train_y[train_y==0]), len(train_y[train_y==1])
        test_1, test_2 = len(test_y[test_y==0]), len(test_y[test_y==1])

        print('{} of KFold {}'.format(i,kfold.n_splits))
        print('                >Train: 1=%d, 2=%d, Test: 1=%d, 2=%d' % (train_1, train_2, test_1, test_2))

        model.fit(train_X, train_y)
        score = roc_auc_score(test_y, model.predict(test_X))

        print('                ROC AUC score:',score)
        cv_score.append(score)

        model_accu_stratified.append(model.score(test_X, test_y))
        i+=1
    print('-----')
    # call to evaluations
    Evaluation(model, model_accu_stratified,cv_score,test_X, test_y, model_name )

```

Evaluation

```

In [10]: def Evaluation(model, model_accu_stratified,cv_score, test_X, test_y, classifier):

    print('List of possible accuracy are :',(round(model_accu_stratified[0], 5))
          , ' ', round(model_accu_stratified[1],5) , ' ', round(model_accu_stratified[2],5))
    print('\nMaximum Accuracy That can be obtained from this model is:',
          max(model_accu_stratified)*100, '%')
    print('\nMinimum Accuracy:',
          min(model_accu_stratified)*100, '%')
    print('\nOverall(Mean) Accuracy:',
          mean(model_accu_stratified)*100, '%')
    print('\nStandard Deviation is:', stdev(model_accu_stratified))
    print('-----')
    print('Cv: ',cv_score,'\nMean cv Score :',np.mean(cv_score))
    print('-----')
    print("\n          Confusion Matrix on tested 5th flod data\n")

    cm =confusion_matrix(test_y,model.predict(test_X))
    print('-----')
    tp, fn, fp, tn =confusion_matrix(test_y,model.predict(test_X)).reshape(-1)

    # classification report for precision, recall f1-score and accuracy
    matrix = classification_report(test_y, model.predict(test_X))
    print('Classification report : \n',matrix)
    print('-----')

    plt.figure(figsize=(9,9))
    sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = 'magma')
    plt.ylabel('Actual label');
    plt.xlabel('Predicted label');
    all_sample_title = 'Accuracy Score is {0} of {1}'.format(round(np.mean(cv_score)*100, 2), 100)
    plt.title(all_sample_title, size = 15);

```

-----Driver Class-----

```
In [16]: load_data()
ExploratoryDataAnalysis(0)
X, y = DataPreprocessing()
load_data().head()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 245057 entries, 0 to 245056
Data columns (total 4 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0    c1      245057 non-null  int64  
 1    c2      245057 non-null  int64  
 2    c3      245057 non-null  int64  
 3    L1      245057 non-null  int64  
dtypes: int64(4)
memory usage: 7.5 MB
```

Out[16]:

	c1	c2	c3	L1
0	74	85	123	1
1	73	84	122	1
2	72	83	121	1
3	70	81	119	1
4	70	81	119	1

1. Decision Tree Classifier

```
In [17]: # Create classifier object.
dtree = DecisionTreeClassifier()
ApplyClassifier(dtree)
```

```
-----
Classifier Name: DecisionTreeClassifier
-----
```

```
-----
1 of KFold 5
```

```
>Train: 1=40686, 2=155358, Test: 1=10172, 2=38840
ROC AUC score: 0.9988706250432824
```

```
2 of KFold 5
```

```
>Train: 1=40686, 2=155359, Test: 1=10172, 2=38839
ROC AUC score: 0.9989689278294233
```

```
3 of KFold 5
```

```
>Train: 1=40686, 2=155359, Test: 1=10172, 2=38839
ROC AUC score: 0.9992638550807013
```

```
4 of KFold 5
```

```
>Train: 1=40687, 2=155358, Test: 1=10171, 2=38840
ROC AUC score: 0.9989384422867423
```

```
5 of KFold 5
```

```
>Train: 1=40687, 2=155358, Test: 1=10171, 2=38840
ROC AUC score: 0.9992614791190113
```

```
-----
List of possible accuracy are : 0.99925 , 0.99929 , 0.99941
```

```
Maximum Accuracy That can be obtained from this model is: 99.94082960967945 %
```

```
Minimum Accuracy: 99.91226459366264 %
```

```
Overall(Mean) Accuracy: 99.92817969636089 %
```

```
Standard Deviation is: 0.00010834921257841509
```


```
-----
Cv: [0.9988706250432824, 0.9989689278294233, 0.9992638550807013, 0.998938442
2867423, 0.9992614791190113]
```

```
Mean cv Score : 0.9990606658718321
-----
```

```
-----
Confusion Matrix on tested 5th fold data
```

```
-----
Classification report :
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10171
1	1.00	1.00	1.00	38840
accuracy			1.00	49011
macro avg	1.00	1.00	1.00	49011
weighted avg	1.00	1.00	1.00	49011



```
In [143]: # Visualising the graph without the use of graphvizplt.figure(figsize = (20,20))
dec_tree = plot_tree(decision_tree=dtree, feature_names = df1.columns,
                    class_names = ["1", "2"], filled = True, precision = 3, rou
```



2. KNeighbors Classifier

In [18]: *# Create classifier object.*

```
k = KNeighborsClassifier()
ApplyClassifier(k)
```

```
-----
Classifier Name: KNeighborsClassifier
-----
```

```
-----
1 of KFold 5
```

```
>Train: 1=40686, 2=155358, Test: 1=10172, 2=38840
ROC AUC score: 0.9996079431714774
```

```
2 of KFold 5
```

```
>Train: 1=40686, 2=155359, Test: 1=10172, 2=38839
ROC AUC score: 0.9996313440998963
```

```
3 of KFold 5
```

```
>Train: 1=40686, 2=155359, Test: 1=10172, 2=38839
ROC AUC score: 0.9997167795257345
```

```
4 of KFold 5
```

```
>Train: 1=40687, 2=155358, Test: 1=10171, 2=38840
ROC AUC score: 0.9995541014359904
```

```
5 of KFold 5
```

```
>Train: 1=40687, 2=155358, Test: 1=10171, 2=38840
ROC AUC score: 0.9996547541165535
-----
```

```
-----
List of possible accuracy are : 0.99955 , 0.99953 , 0.99955
```

```
Maximum Accuracy That can be obtained from this model is: 99.9551130335428 %
```

```
Minimum Accuracy: 99.94082960967945 %
```

```
Overall(Mean) Accuracy: 99.95103158428502 %
```

```
Standard Deviation is: 5.948772589477762e-05
-----
```


```
-----
Cv: [0.9996079431714774, 0.9996313440998963, 0.9997167795257345, 0.999554101
4359904, 0.9996547541165535]
```

```
Mean cv Score : 0.9996329844699303
-----
```

```
-----
Confusion Matrix on tested 5th flod data
-----
```

```
-----
Classification report :
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10171
1	1.00	1.00	1.00	38840
accuracy			1.00	49011
macro avg	1.00	1.00	1.00	49011
weighted avg	1.00	1.00	1.00	49011



3. AdaBoostClassifier

```
In [19]: # Create classifier object.
Ada = AdaBoostClassifier(n_estimators=50, learning_rate=1.0, algorithm='SAMME.R')
ApplyClassifier(Ada)
```

```
-----
Classifier Name: AdaBoostClassifier
-----
```

```
-----
1 of KFold 5
```

```
>Train: 1=40686, 2=155358, Test: 1=10172, 2=38840
ROC AUC score: 0.9235979362989536
```

```
2 of KFold 5
```

```
>Train: 1=40686, 2=155359, Test: 1=10172, 2=38839
ROC AUC score: 0.9227771870418569
```

```
3 of KFold 5
```

```
>Train: 1=40686, 2=155359, Test: 1=10172, 2=38839
ROC AUC score: 0.9245092964566701
```

```
4 of KFold 5
```

```
>Train: 1=40687, 2=155358, Test: 1=10171, 2=38840
ROC AUC score: 0.9282899038187469
```

```
5 of KFold 5
```

```
>Train: 1=40687, 2=155358, Test: 1=10171, 2=38840
ROC AUC score: 0.9231340410089429
```

```
-----
List of possible accuracy are : 0.95332 , 0.95368 , 0.9547
```

```
Maximum Accuracy That can be obtained from this model is: 95.53569606822958 %
```

```
Minimum Accuracy: 95.30513558180816 %
```

```
Overall(Mean) Accuracy: 95.40227568604807 %
```

```
Standard Deviation is: 0.0009746834098173316
```


```
-----
Cv: [0.9235979362989536, 0.9227771870418569, 0.9245092964566701, 0.928289903
8187469, 0.9231340410089429]
```

```
Mean cv Score : 0.9244616729250341
-----
```

```
-----
Confusion Matrix on tested 5th fold data
```

```
-----
Classification report :
```

	precision	recall	f1-score	support
0	0.90	0.87	0.89	10171
1	0.97	0.97	0.97	38840
accuracy			0.95	49011
macro avg	0.93	0.92	0.93	49011
weighted avg	0.95	0.95	0.95	49011



4. LogisticRegression

In [20]: *# Create classifier object.*

```
lr = LogisticRegression()
ApplyClassifier(lr)
```

```
-----
Classifier Name: LogisticRegression
-----
```

```
-----
1 of KFold 5
```

```
>Train: 1=40686, 2=155358, Test: 1=10172, 2=38840
ROC AUC score: 0.8799387355204185
```

```
2 of KFold 5
```

```
>Train: 1=40686, 2=155359, Test: 1=10172, 2=38839
ROC AUC score: 0.8817356745017649
```

```
3 of KFold 5
```

```
>Train: 1=40686, 2=155359, Test: 1=10172, 2=38839
ROC AUC score: 0.8841933952677609
```

```
4 of KFold 5
```

```
>Train: 1=40687, 2=155358, Test: 1=10171, 2=38840
ROC AUC score: 0.8877354017667605
```

```
5 of KFold 5
```

```
>Train: 1=40687, 2=155358, Test: 1=10171, 2=38840
ROC AUC score: 0.8844174009099395
```

```
-----
List of possible accuracy are : 0.91753 , 0.91773 , 0.91939
```

```
Maximum Accuracy That can be obtained from this model is: 92.14462059537655 %
```

```
Minimum Accuracy: 91.75304007181914 %
```

```
Overall(Mean) Accuracy: 91.87899959992626 %
```

```
Standard Deviation is: 0.001657425755787211
```

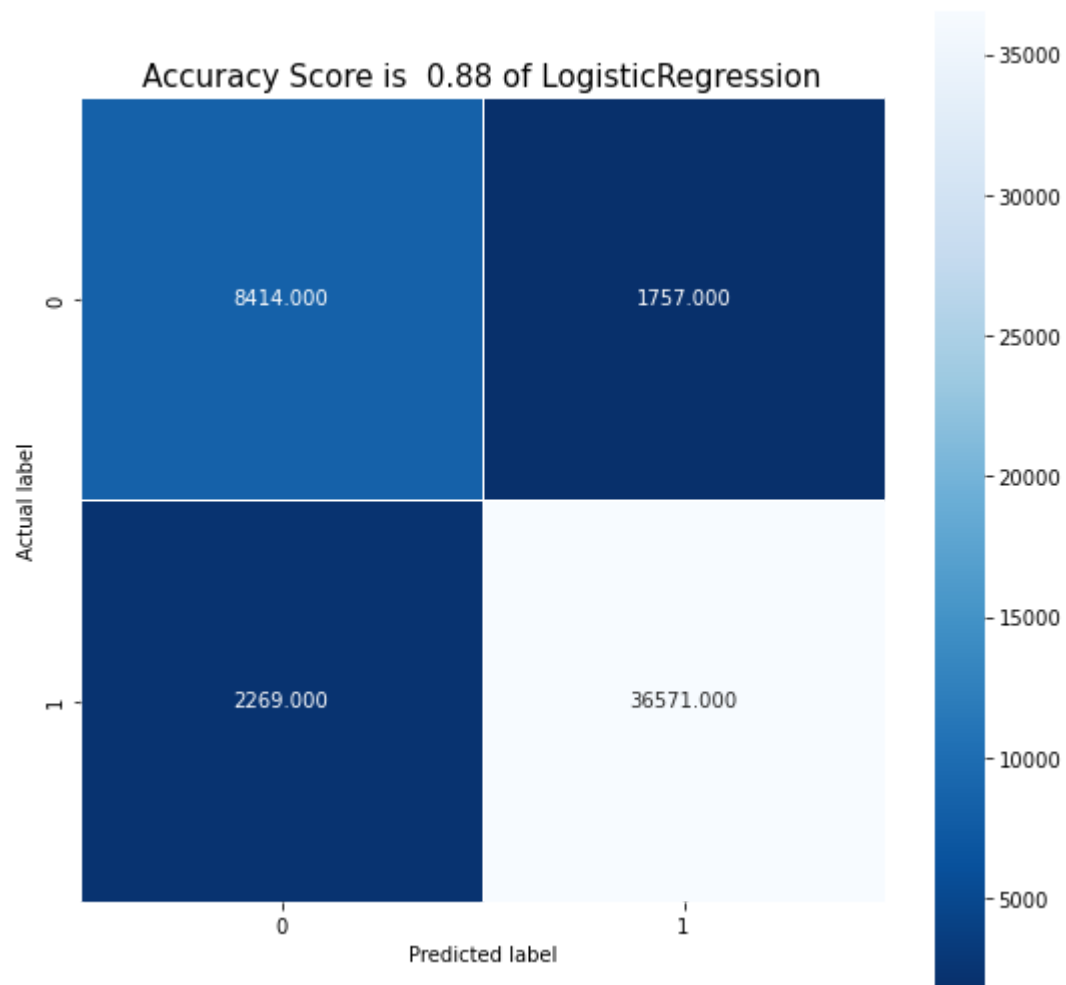
```
-----
Cv: [0.8799387355204185, 0.8817356745017649, 0.8841933952677609, 0.887735401
7667605, 0.8844174009099395]
```

```
Mean cv Score : 0.8836041215933289
-----
```

```
-----
Confusion Matrix on tested 5th fold data
```

```
-----
Classification report :
```

	precision	recall	f1-score	support
0	0.79	0.83	0.81	10171
1	0.95	0.94	0.95	38840
accuracy			0.92	49011
macro avg	0.87	0.88	0.88	49011
weighted avg	0.92	0.92	0.92	49011



5. Naive Bayes Classification

In [21]: *# Create classifier object.*

```
nBayes = GaussianNB()
ApplyClassifier(nBayes)
```

```
-----
Classifier Name: GaussianNB
-----
```

```
-----
1 of KFold 5
```

```
>Train: 1=40686, 2=155358, Test: 1=10172, 2=38840
ROC AUC score: 0.8494471607405154
```

```
2 of KFold 5
```

```
>Train: 1=40686, 2=155359, Test: 1=10172, 2=38839
ROC AUC score: 0.8542522613468588
```

```
3 of KFold 5
```

```
>Train: 1=40686, 2=155359, Test: 1=10172, 2=38839
ROC AUC score: 0.855135863563809
```

```
4 of KFold 5
```

```
>Train: 1=40687, 2=155358, Test: 1=10171, 2=38840
ROC AUC score: 0.8577942302487404
```

```
5 of KFold 5
```

```
>Train: 1=40687, 2=155358, Test: 1=10171, 2=38840
ROC AUC score: 0.8553469477293583
-----
```

```
-----
List of possible accuracy are : 0.92188 , 0.92357 , 0.92485
```

```
Maximum Accuracy That can be obtained from this model is: 92.57513619391565 %
```

```
Minimum Accuracy: 92.18762751979106 %
```

```
Overall(Mean) Accuracy: 92.38949546988422 %
```

```
Standard Deviation is: 0.0014805093501863067
-----
```


```
-----
Cv: [0.8494471607405154, 0.8542522613468588, 0.855135863563809, 0.8577942302
487404, 0.8553469477293583]
```

```
Mean cv Score : 0.8543952927258565
-----
```

```
-----
Confusion Matrix on tested 5th fold data
-----
```

```
-----
Classification report :
```

	precision	recall	f1-score	support
0	0.87	0.74	0.80	10171
1	0.93	0.97	0.95	38840
accuracy			0.92	49011
macro avg	0.90	0.86	0.88	49011
weighted avg	0.92	0.92	0.92	49011



In []:

